




A Held-out prediction across training budgets



B Improvement per pass versus the plateau rule



C Held-out minus training likelihood



D Learning rate and optimiser sensitivity

